

---

# Advanced Artificial Intelligence

## Lab 06

# Outline

---

- More details on the content in lecture 06
- Exercise

# Details

---

- **Linear Regression Example**
- Cross Validation Example
- Gradient Descent

# Sklearn library




The sklearn library contains a lot of efficient tools for machine learning modeling including classification, regression, clustering and dimensionality reduction

## Classification

Identifying which category an object belongs to.

**Applications:** Spam detection, image recognition.

**Algorithms:** SVM, nearest neighbors, random forest, and more...




## Regression

Predicting a continuous-valued attribute associated with an object.

**Applications:** Drug response, Stock prices.

**Algorithms:** SVR, nearest neighbors, random forest, and more...



# Diabetes dataset

<b>Number of Instances:</b>	442
<b>Number of Attributes:</b>	First 10 columns are numeric predictive values
<b>Target:</b>	Column 11 is a quantitative measure of disease progression one year after baseline
<b>Attribute Information:</b>	<ul style="list-style-type: none"><li>• age age in years</li><li>• sex</li><li>• bmi body mass index</li><li>• bp average blood pressure</li><li>• s1 tc, total serum cholesterol</li><li>• s2 ldl, low-density lipoproteins</li><li>• s3 hdl, high-density lipoproteins</li><li>• s4 tch, total cholesterol / HDL</li><li>• s5 ltg, possibly log of serum triglycerides level</li><li>• s6 glu, blood sugar level</li></ul>

# Linear Model


---

$$\hat{y}(w, x) = w_0 + w_1x_1 + \dots + w_px_p$$

we designate the vector  $w = (w_1, \dots, w_p)$  as `coef_` and  $w_0$  as `intercept_`.

# Linear Regression

---



[LinearRegression](#) fits a linear model with coefficients  $w = (w_1, \dots, w_p)$  to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation. Mathematically it solves a problem of the form:

$$\min_w ||Xw - y||_2^2$$

```
from sklearn import linear_model
l_reg = linear_model.LinearRegression()

X1 = [[0,0],[1,1],[2,2]]
y1 = [0,1,2]

l_reg.fit(X1,y1)
print(l_reg.coef_,l_reg.intercept_)

✓

[0.5 0.5] 2.220446049250313e-16
```

# Load the diabetes dataset

---

```
from sklearn import datasets
diabetes_X, diabetes_y = datasets.load_diabetes(return_X_y=True)
print(diabetes_X.shape, diabetes_y.shape)
print(diabetes_X[:3])
print(diabetes_y[:3])
```

✓

```
(442, 10) (442,)
[[ 0.03807591  0.05068012  0.06169621  0.02187235 -0.0442235  -0.03482076
 -0.04340085 -0.00259226  0.01990842 -0.01764613]
 [-0.00188202 -0.04464164 -0.05147406 -0.02632783 -0.00844872 -0.01916334
  0.07441156 -0.03949338 -0.06832974 -0.09220405]
 [ 0.08529891  0.05068012  0.04445121 -0.00567061 -0.04559945 -0.03419447
 -0.03235593 -0.00259226  0.00286377 -0.02593034]]
[151.  75. 141.]
```




# Split the diabetes dataset

```
import numpy as np
# Use only one feature
diabetes_X = diabetes_X[:, np.newaxis, 2]

# Split the data into training/testing sets
diabetes_X_train = diabetes_X[:-20]
diabetes_X_test = diabetes_X[-20:]

# Split the targets into training/testing sets
diabetes_y_train = diabetes_y[:-20]
diabetes_y_test = diabetes_y[-20:]
```



✓

```
import numpy as np
print(diabetes_X[:3])
print('*****')
print(diabetes_X[:3,2])
print('*****')
print(diabetes_X[:3, np.newaxis, 2])
```

✓

```
[[ 0.03807591  0.05068012  0.06169621  0.02187235 -0.0442235 -0.03482076
 -0.04340085 -0.00259226  0.01990842 -0.01764613]
 [-0.00188202 -0.04464164 -0.05147406 -0.02632783 -0.00844872 -0.01916334
  0.07441156 -0.03949338 -0.06832974 -0.09220405]
 [ 0.08529891  0.05068012  0.04445121 -0.00567061 -0.04559945 -0.03419447
 -0.03235593 -0.00259226  0.00286377 -0.02593034]]
```

\*\*\*\*\*

```
[ 0.06169621 -0.05147406  0.04445121]
```

\*\*\*\*\*

```
[[ 0.06169621]
 [-0.05147406]
 [ 0.04445121]]
```

# Train the model and predict

---

```
from sklearn import linear_model
from sklearn.metrics import mean_squared_error

# Create linear regression object
regr = linear_model.LinearRegression()

# Train the model using the training sets
regr.fit(diabetes_X_train, diabetes_y_train)

# Make predictions using the testing set
diabetes_y_pred = regr.predict(diabetes_X_test)


print('Mean squared error: %.2f'
      % mean_squared_error(diabetes_y_test, diabetes_y_pred))
```

✓

Mean squared error: 2548.07

# 1-Dimension feature

---



We only use the **one feature** of the diabetes dataset, in order to illustrate the data points within the two-dimensional plot.

[Linear Regression Example — scikit-learn 1.0.1 documentation](#)


# Visualize the model

---

```
print('Coefficients:', regr.coef_, 'intercept:', regr.intercept_)  
import matplotlib.pyplot as plt  
plt.scatter(diabetes_X_test, diabetes_y_test, color='black')  
plt.plot(diabetes_X_test, diabetes_y_pred, color='blue', linewidth=3)  
  
plt.xticks(())  
plt.yticks(())  
  
plt.show()
```

✓

Coefficients: [938.23786125] intercept: 152.91886182616167



# Details

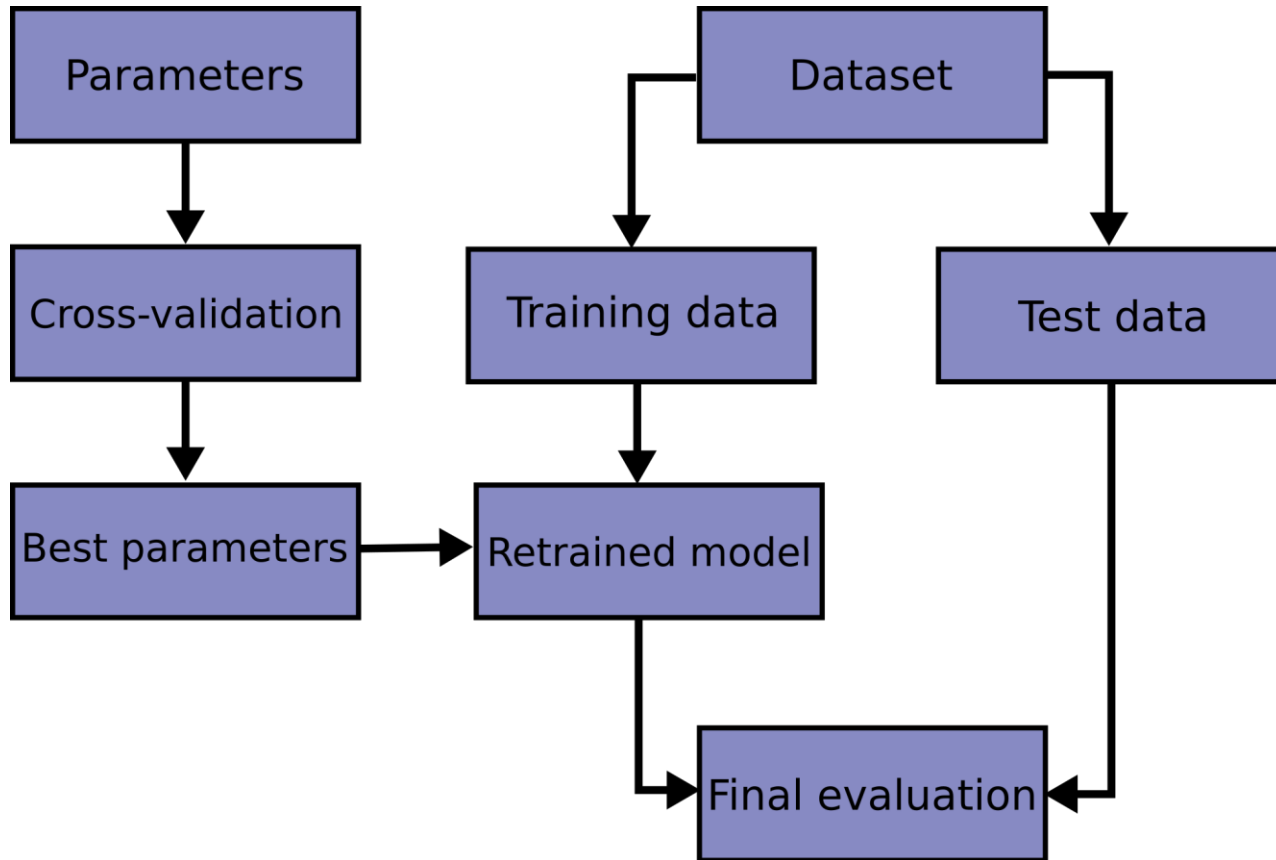
---

- Linear Regression Example
- **Cross Validation Example**
- Gradient Descent

# Cross Validation

---

We use CV to select best  
hyperparameters



# Load the dataset

---

Let's load the iris data set to fit a linear support vector machine on it

```
>>> import numpy as np
>>> from sklearn.model_selection import train_test_split
>>> from sklearn import datasets
>>> from sklearn import svm

>>> X, y = datasets.load_iris(return_X_y=True)
>>> X.shape, y.shape
((150, 4), (150,))
```

# Split Dataset, Train and Test

---

```
>>> X_train, X_test, y_train, y_test = train_test_split(
...     X, y, test_size=0.4, random_state=0)

>>> X_train.shape, y_train.shape
((90, 4), (90,))
>>> X_test.shape, y_test.shape
((60, 4), (60,))


>>> clf = svm.SVC(kernel='linear', C=1).fit(X_train, y_train)
>>> clf.score(X_test, y_test)
0.96...
```

We can now quickly sample a training set while holding out 40% of the data for testing (evaluating) our classifier



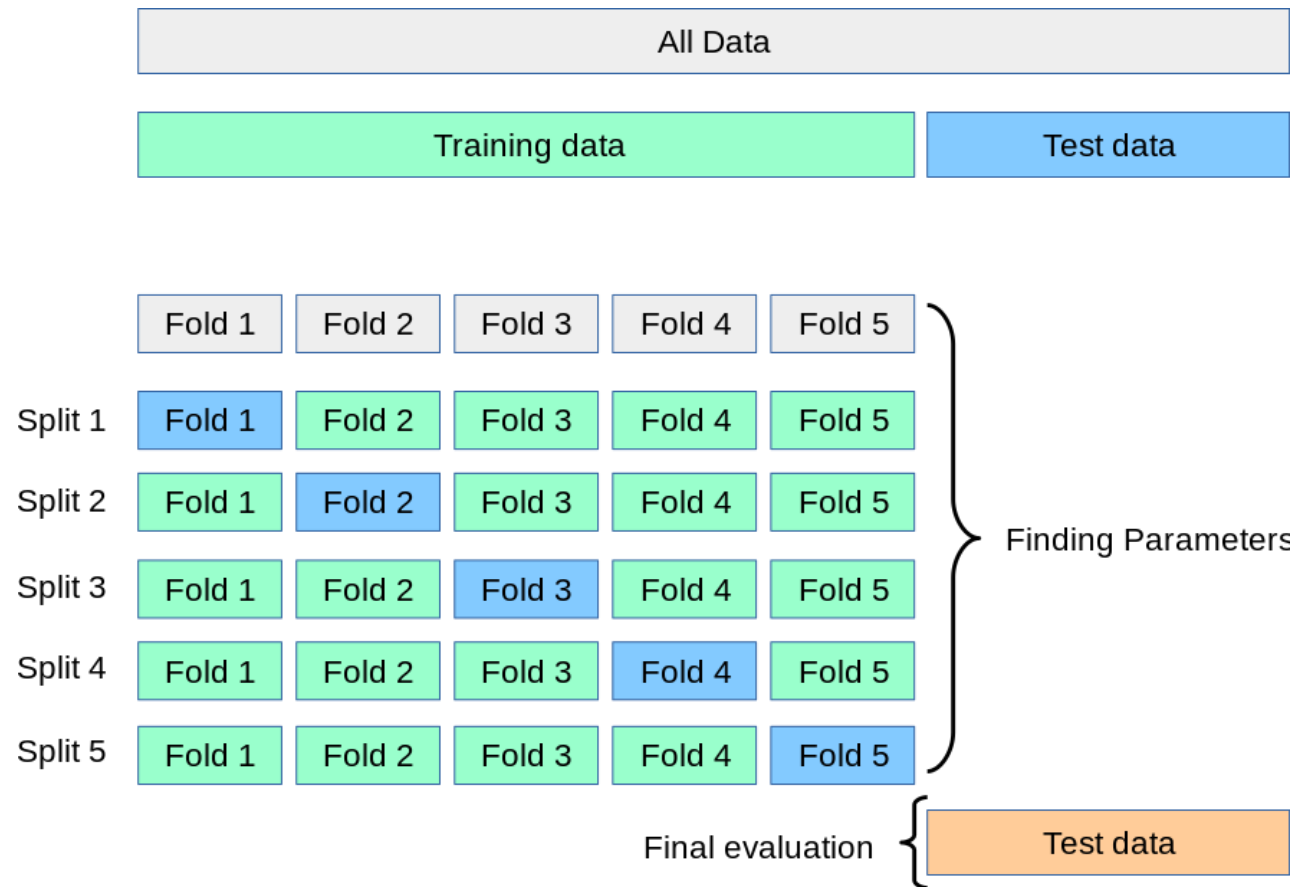
# Validation Set

---



By partitioning the available data into three sets, we drastically reduce the number of samples which can be used for learning the model

# K-Fold cross validation



- A model is trained using  $k-1$  of the folds as training data;
- the resulting model is validated on the remaining part of the data

# 5-fold CV

---

```
>>> from sklearn.model_selection import cross_val_score
>>> clf = svm.SVC(kernel='linear', C=1, random_state=42)
>>> scores = cross_val_score(clf, X, y, cv=5)
>>> scores
array([0.96..., 1. , 0.96..., 0.96..., 1. ])
```

Estimate the accuracy by splitting the data, fitting a model, and computing the score for 5 consecutive times (with different splits each time)

# Leave One Out (LOO)

---

```
>>> from sklearn.model_selection import LeaveOneOut

>>> X = [1, 2, 3, 4]
>>> loo = LeaveOneOut()
>>> for train, test in loo.split(X):
...     print("%s %s" % (train, test))
[1 2 3] [0]
[0 2 3] [1]
[0 1 3] [2]
[0 1 2] [3]
```


Each learning set is created by taking all the samples except one, the test set being the sample left out.

# Details

---

- Linear Regression Example
- Cross Validation Example
- **Gradient Descent**

$$f(x) = x * x$$



$$f(x) = x^2 \quad f'(x) = 2x \quad \text{step} = -0.1$$

```
def f(x):  
    return x*2  
  
def df(x):  
    return 2*x  
  
step = -0.1  
  
x = 3  
  
for i in range(20):  
    print(x)  
    x = df(x)*step+x
```



```
3  
2.4  
1.92  
1.536  
1.2288000000000001  
0.9830400000000001  
0.7864320000000001  
0.6291456000000001  
0.5033164800000001  
0.40265318400000005  
0.32212254720000005  
0.25769803776000005  
0.20615843020800004  
0.16492674416640002  
0.13194139533312002  
0.10555311626649602  
0.08444249301319681  
0.06755399441055746  
0.05404319552844596  
0.04323455642275677
```

$$f(x, y) = 2x^2 + y^2$$

$$\frac{\partial f}{\partial x} = 4x$$

$$\frac{\partial f}{\partial y} = 2y$$

$$\nabla f(x, y) = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]^T$$

$$\nabla f(x, y) = [4x, 2y]^T$$

```
def f(x,y):  
    return 2*x*x+y*y
```

```
def dx(x):  
    return 4*x
```

```
def dy(y):  
    return 2*y
```

```
step = -0.1
```

```
x = 3  
y = 5
```

```
for i in range(20):  
    print(f(x,y))  
    x = dx(x)*step+x  
    y = dy(y)*step+y
```

✓

```
43  
22.479999999999997  
12.5728  
7.393408  
4.496634879999999  
2.7931936768  
1.7571690004480003  
1.1136171773132804  
0.7087654396100611  
0.45218804195708123  
0.2888884846709232  
0.1847043598040117  
0.11814445293583098  
0.07558856843698666  
0.04836808648057122  
0.03095248031268949  
0.019808473187565876  
0.012677021723522221  
0.008113149501107041  
0.005192363696007523
```

# n-variable function

---

In the same way, if we get a function with 4 variables, we would get a gradient vector with 4 partial derivatives.

Generally, an  $n$ -variable function results in an  $n$ -dimensional gradient vector.

$$f(x_1, x_2, \dots, x_n) = \nabla f = \left[ \frac{\partial}{\partial x_1}, \frac{\partial}{\partial x_2}, \dots, \frac{\partial}{\partial x_n} \right]$$



# Performance measure

---

$$Err = \frac{1}{N} \sum_{\mu} \overbrace{(y^{\mu} - (W_1 X^{\mu} + W_0))}^{\text{Ground truth}}^2$$

Predicted values

---

**MSE**

# Partial derivatives

---

$$Error = \frac{1}{N} \sum_{\mu} (\overset{\text{Ground truth}}{y^{\mu}} - \underbrace{(W_1 X^{\mu} + W_0)}_{\text{Predicted values}})^2$$

---

**MSE**

$$\frac{\partial}{\partial W_0} = -\frac{2}{N} \sum_{\mu} (y^{\mu} - (W_1 X^{\mu} + W_0))$$

$$\frac{\partial}{\partial W_1} = -\frac{2}{N} \sum_{\mu} X \times (y^{\mu} - (W_1 \times X^{\mu} + W_0))$$

# The gradient

---

$$\nabla Err = \left[ \frac{\partial}{\partial W_0}, \frac{\partial}{\partial W_1} \right]^T$$

$$W_0 = W_0 - \eta \left( \frac{\partial}{\partial W_0} \right)$$

$$W_1 = W_1 - \eta \left( \frac{\partial}{\partial W_1} \right)$$

# The gradient

---

```
for epoch in range(total_epochs):
    for X_batch, y_batch in next_batch(Xs_train, Ys_train, batch_size=batch_size):

        # linearly combine input and weights
        train_pred = W0 + np.dot(X_batch, W1)


        # calculate the SSE between predicted and true values
        train_err = mean_squared_error(y_batch, train_pred)

        # calculate the gradients with respect to W0 and W1
        DW0 = -(2/batch_size) * sum(y_batch.squeeze() - train_pred.squeeze())
        DW1 = -(2/batch_size) * sum(X_batch.squeeze() * (y_batch.squeeze() - train_pred.squeeze()))

        # update W0 and W1 in the opposite direction to the gradient
        W0 = W0 - lr * DW0
        W1 = W1 - lr * DW1
```

# The gradient

---



# Exercise

---

1. Compare linear regression and SVM using the sklearn library
2. Try to implement linear regression from scratch with Python