
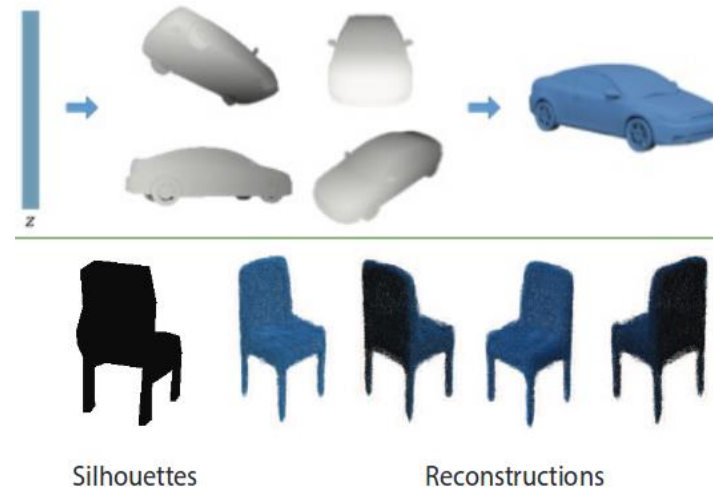

Feature Engineering



“Sense” of Human

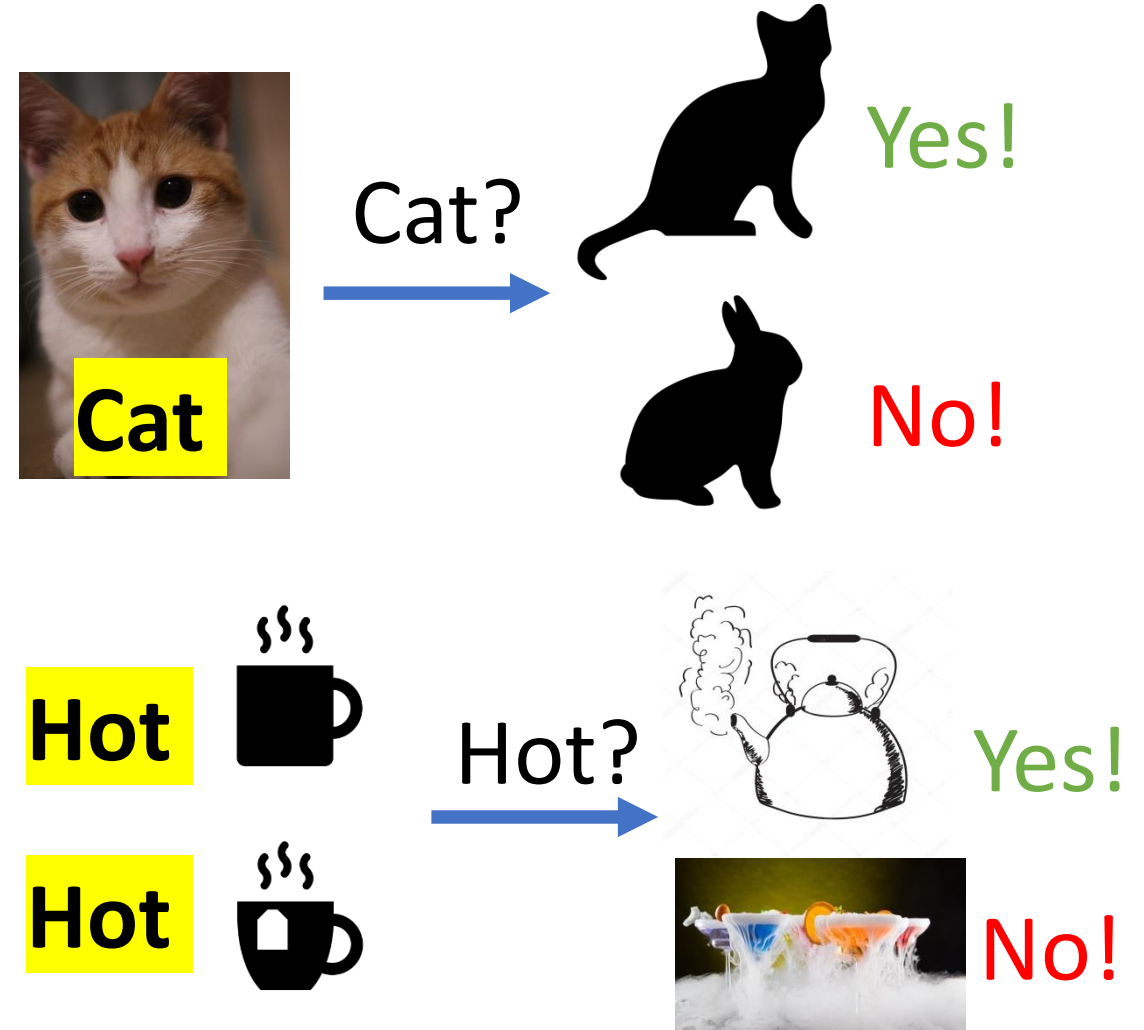

- Definition of "**sense**":
 - "A system that consists of a group of sensory cell types that responds to a specific physical phenomenon, and that corresponds to a particular group of regions within the brain where the signals are **received** and **interpreted**."
 - "A physiological capacity of organisms that provides data for **perception**."



By Arsi Warrior - Own work, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=68328908>

Five “Traditional” Senses of a Human

- Sight
- Hearing
- Taste
- Smell
- Touch



What is the Perception of Machines?

- How to represent the world to a machine?
 - How can you interpret “hot”/“soft”?
 - How can you represent the chair looked from different directions?

=> Mapping raw data to features.
- How to select the sensory information to be sent to a machine?
 - How to decide which information is more important?
 - How to extract this information and transform the data correctly?
- How to learn from experience/examples?
 - Again, **generalization**: the ability to categorize correctly new examples that differ from those used for training.

Why Feature Engineering?

- Help the model to understand the data set as the same or similar way the human beings do.
- If the quality and size of the data are terrible, training longer or using a deeper network won't help.
- The pre-processing of data and feature engineering are the foundation of the pyramid.
- Preparing a better dataset can be more important than tuning the parameters for your model.

Representation of The Real World

- AI provides human with powerful tools for ‘better’ decision making.
- To accomplish a task, AI needs human to :
 - Formulate the real-world problem to those that can be read by computers.
 - Choose a model of the task & choose learning algorithms for the model.
 - Find useful raw data of the task.
 - Convert raw data to the formats that the computer can read → e.g. features.

Outline

- Features and Feature Engineering
- Tackling Feature Explosion

Features and Feature Engineering

Introduction

Transforming Data


Examples

Feature

- **Feature**: information that describes a problem at hand and is potentially useful for prediction / problem-solving.

Feature Engineering

- **Feature engineering**: the process of determining which features might be useful in training a model, then creating these features by transforming raw data.
- In short: design and process features for AI applications.
 - An informal terminology, but is considered essential in applied AI.



Feature Learning / Extraction

- Feature learning process
 - 1) understand the properties of the task and how they may interact with the strengths and limitations of the chosen model
 - 2) design a set of features
 - 3) run experiments and analyze the results on a validation dataset
 - 4) change the feature set
 - 5) go to 2).
- Difficult and expensive.
- Automated feature learning is preferred.

Features and Feature Engineering

Introduction

Transforming Data

Examples

Feature Types

- Numerical features
 - Floats
 - Integers
- Categorical features
 - Discrete set of possible values (e.g., names of students of AAI): One-/Multi-hot encoding to map the categorical data to binary vectors.
- Image features.
- ...

Need to transform data!

Reasons for Data Transformation [8]

- **Mandatory transformations** for data compatibility.
 - Converting non-numeric features into numeric.
 - Resize inputs to a fixed size.
- **Optional quality transformations:** may help the model perform better.
 - Tokenization or lower-casing of text features.
 - Normalized numeric features.
 - Allowing linear models to introduce non-linearities into the feature space.

Why Normalize Numeric Features?

- **Normalization is necessary**
 - If you have **very different values within the same feature**.
 - Without normalization, your training could blow up with *NaNs* if the gradient update is too large..
 - If you have **two different features with widely different ranges**.
 - This may cause the gradient descent to "bounce" and slow down convergence.
 - A possible solution: using **heterogeneous learning rate**.

Normalization Techniques

Normalization Technique	Formula	When to Use
Linear Scaling	$x' = (x - x_{min}) / (x_{max} - x_{min})$	When the feature is more-or-less uniformly distributed across a fixed range.
Clipping	if $x > \max$, then $x' = \max$. if $x < \min$, then $x' = \min$	When the feature contains some extreme outliers.
Log Scaling	$x' = \log(x)$	When the feature conforms to the power law.
Z-score	$x' = (x - \mu) / \sigma$	When the feature distribution does not contain extreme outliers.

Table source: <https://developers.google.com/machine-learning/data-prep/transform/normalization>

Bucketing

- Sometimes, you need to transform numeric features into categorical features, using a set of thresholds.
=> Bucketing.
- Quantization.

Transforming Categorical Data

- One-/Multi-hot encoding
- Hashing
- Embeddings: A categorical feature represented as a continuous-valued feature (high-dimensional vector -> low-dimensional space).

One-/Multi-hot encoding: Example

One-hot encoding

1. "I have a cat."
2. "Cats have fur."

	I	cat	a	have	fur
I	1	0	0	0	0
have	0	0	0	1	0
a	0	0	1	0	0
cat	0	1	0	0	0

	I	cat	a	have	fur
cat	0	1	0	0	0
have	0	0	0	1	0
fur	0	0	0	0	1

	I	cat	a	have	fur
1	1	1	1	1	0
2	0	1	0	1	1

← Multi-hot encoding

Features and Feature Engineering

Introduction

Transforming Data

Examples

(1) Numerical Features

What you see



What a computer see (raw data)

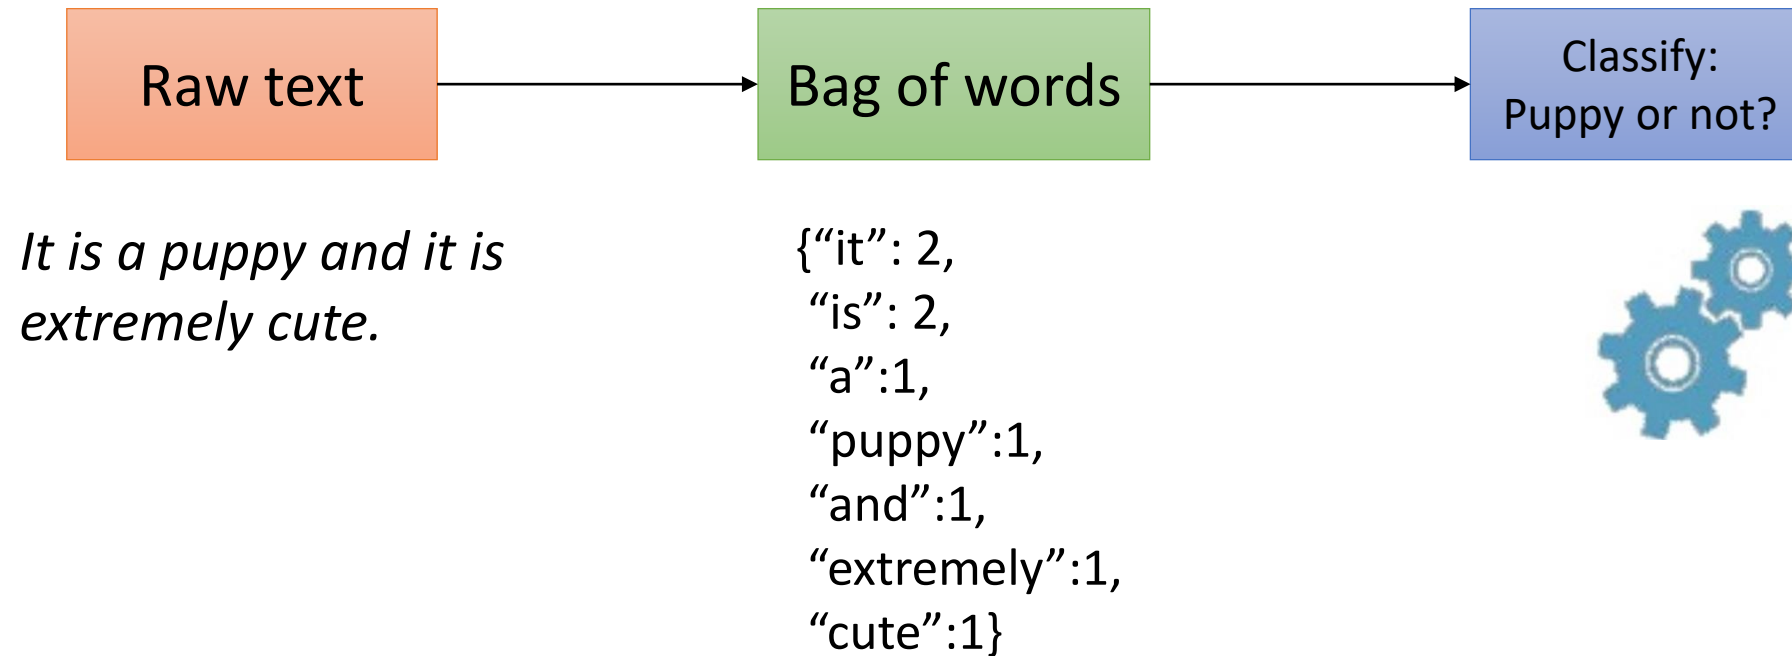
```
StateObservation{  
  gameScore=0,  
  gameTick=0,  
  gameWinner=NO_WINNER,  
  isGameOver=false,  
  worldDimension=[250.0, 200.0],  
  blockSize=10,  
  noOfPlayers=1,  
  ...  
}
```

$[0.0, 0.0, 0.0, -1.0, 250.0, \dots]$

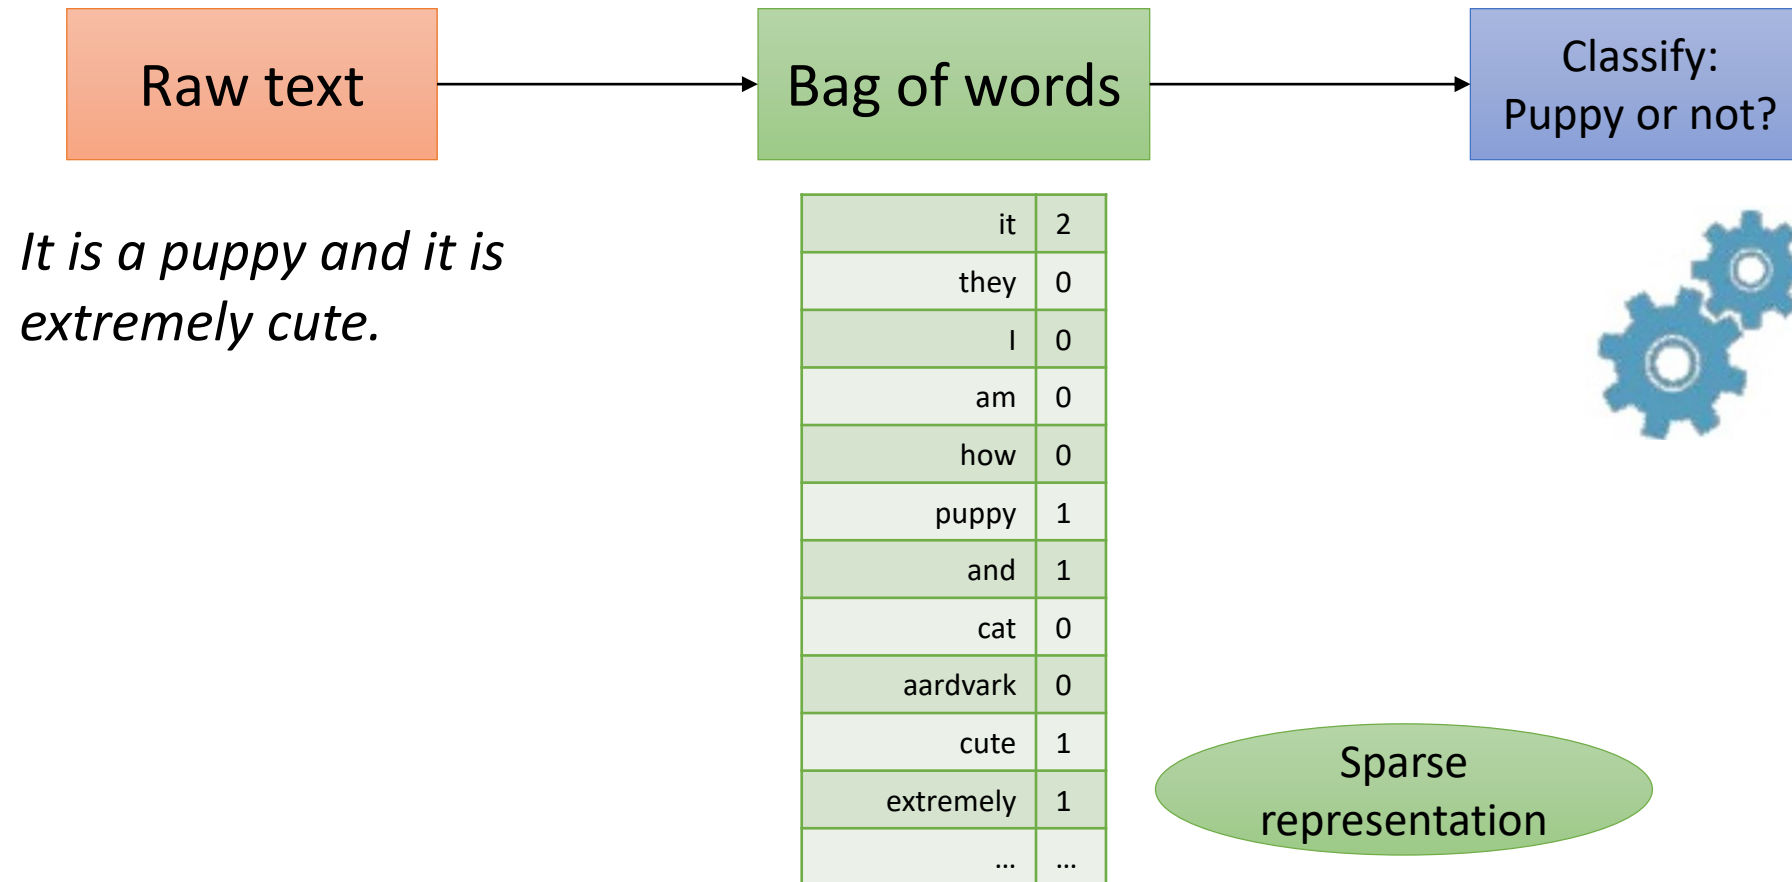
Feature vector

Feature Engineering

(2) Text Features



(2) Text Features (*Continued*)



(3) Image Features



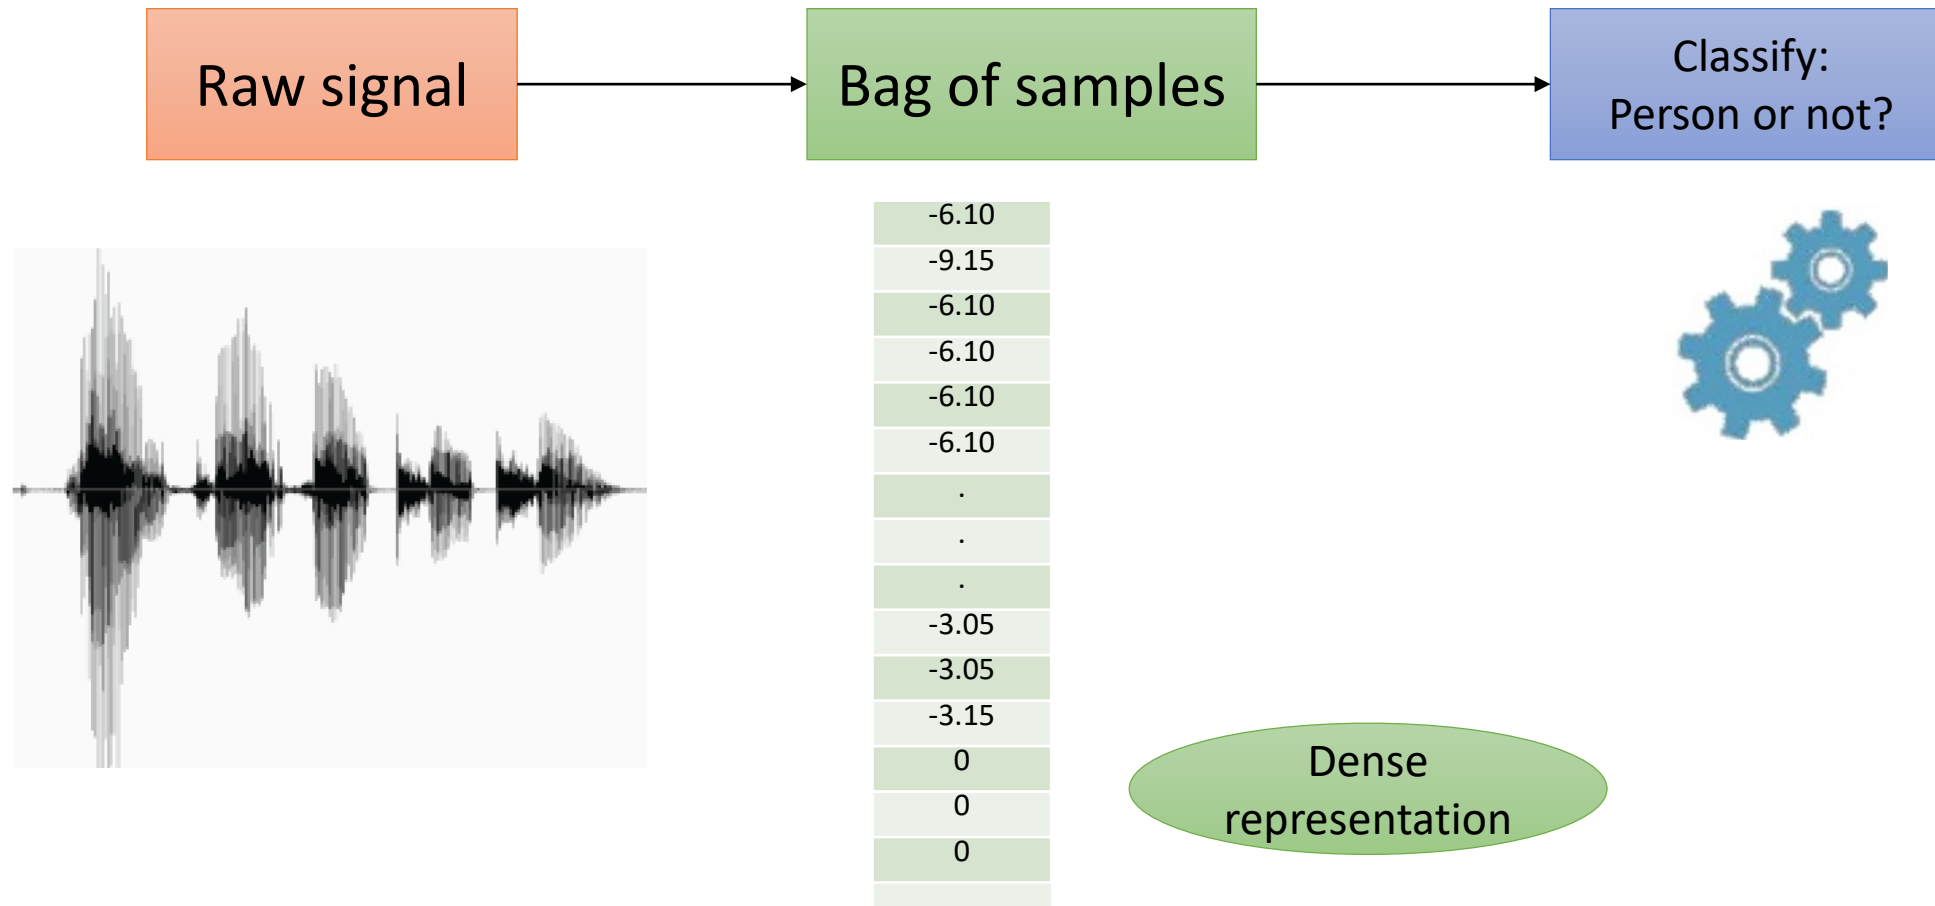


Image source: "Recognizing and learning object categories,"
Li Fei-Fei, Rob Fergus, Anthony Torralba, ICCV 2005—2009.

(3) Image Features



(4) Signal Features



Tackling Feature Explosion

- ❖ Introduction to Feature Explosion
- ❖ Feature Selection
- ❖ Regularization (more detailed)

Feature Explosion

- **Initial features** are always an expression of prior knowledge.
 - text: words, grammatical classes and relations, etc.
 - image: pixels, contours, textures, etc.
 - signal: samples, spectrograms, etc.
- **Feature combinations** might work better.
- Both lead to (extremely) large number of features.
- Too many features become a problem given the limited size of training data. → **overfitting**.

Problems of Feature Explosion

- Storage cost
- Irrelevant, redundant or even harmful features
- Large number of required training samples
 - Adding another feature need **exponential** increase in training samples.
- Dysfunctional distance functions
 - When a measure such as Euclidean distance is used, there is little difference in distance between different pairs of samples.

Benefits of Small Feature Set

- Lead to simpler models.
- Easier to interpret by researchers/users.
- Shorter training times.
- Less computational burden.
- Enhanced generalization by reducing overfitting.
- Reduced feature measurement cost.
- ...

Dealing with Feature Explosion

- **Feature selection:** could use a greedy method.
 - Select some of the features that can reach some best 'criterion'.
- **Regularization:**
 - Include all possible features.
 - Penalize 'complex' hypothesis.

Tackling Feature Explosion

- ❖ Introduction to Feature Explosion
- ❖ Feature Selection
- ❖ Regularization (more detailed)

Selecting Feature Subset

- Reduce the original feature space by throwing out some features.
- **Assumption**: features are redundant or irrelevant.
- **Motivation**: Training data are limited.
 - Restricting #features is a feasible control mechanism.
 - Compact and representative explanation of the task follows Occam's razor.
- **Research Question**: How to select 'good' features from the feature space?
- Feature selection is a **search problem**.

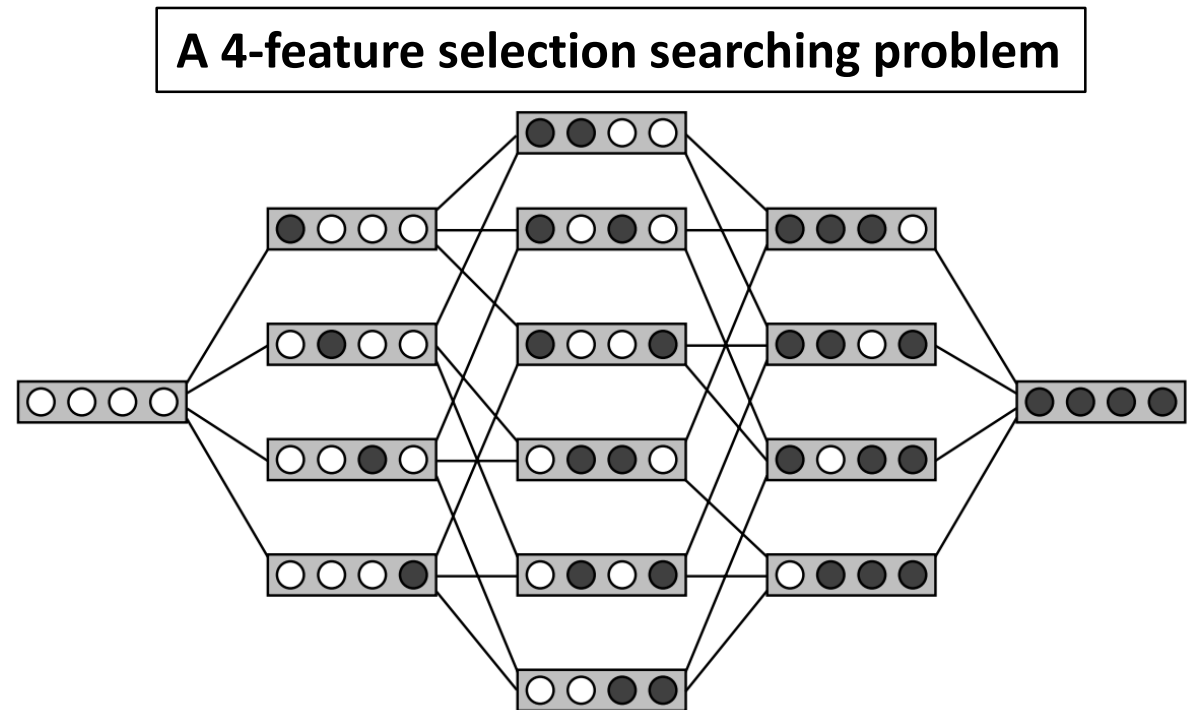
Feature Selection is a Search Problem

- **The state-space formulation:**

- **states:** all possible feature subset
- **initial state:** ?
- **actions:** ?
- **next state:** updated feature subset
- **goal test:** ?
- **cost:** computational cost

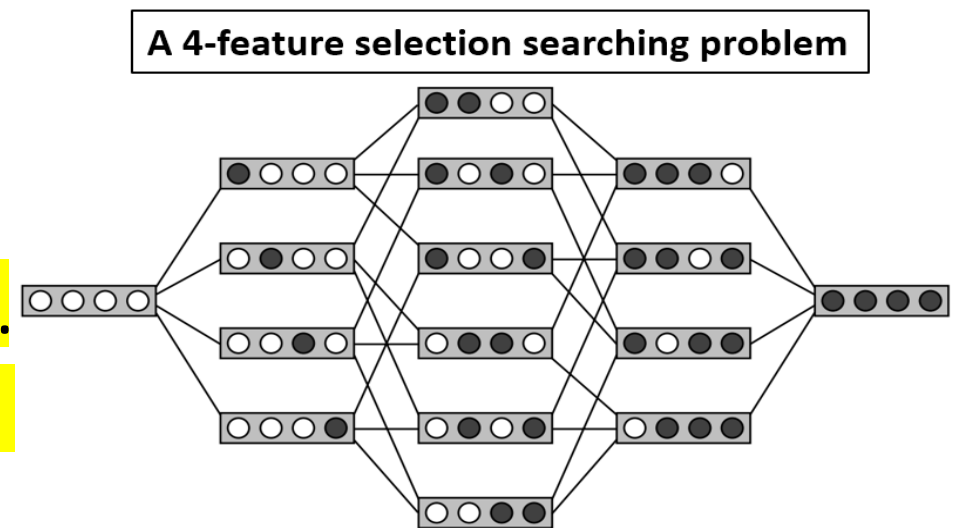
- **Technical question:**

- How to search?
- How to evaluate selected features?
- When to stop?



Search Space of Feature Selection

- **Question:** How large is the search space?
- **Answer:** 2^d , d is #features.
- The search space of the illustration is $2^4 = 16$ and is **feasible to search**.
- When d gets larger and larger, it will become **infeasible to search in practice**.
- We need **heuristics** to guide our search
→ **heuristic search**.



Heuristic Search for Feature Selection

- **Question:** How to do heuristic search in the entire 2^d space?
- **One possible idea:** Greedy heuristic search.

Initial State

- **Empty feature set:** one starts with an empty set and progressively add features yielding to the improvement of a performance index.
→ **forward selection.**
- **Full feature set:** one starts with all the features and progressively eliminate the least useful ones.
→ **backward elimination.**

Actions

- **Forward selection:** add one feature each step.
- **Backward elimination:** remove one feature each step.
- They are called *sequential feature selection* (SFS) methods.
- **Optimality:** They do not examine all possible feature subsets, so no guarantee of finding the optimal subset.

Compare Forward Selection to Backward Elimination

- Both procedures are reasonably fast and robust against overfitting
- Both procedures provide **nested** feature subsets.
- However, they may lead to **different** subsets and one may be **preferred** over the other.

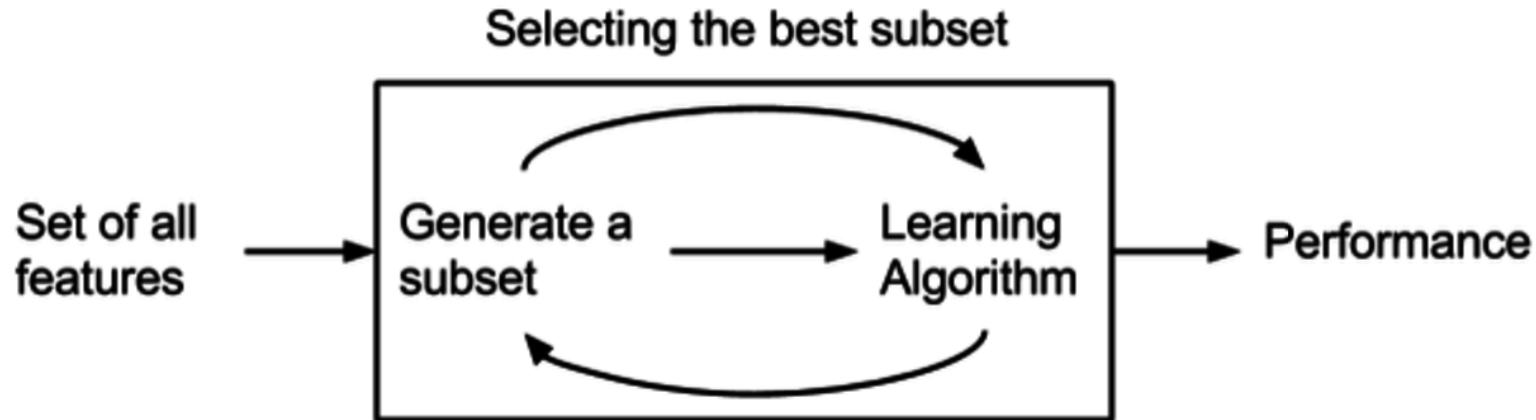
Goal Test

- How to evaluate selected features? e.g.,
 - information theory;
 - prediction accuracy on the training set or validation set.
- When to stop?
 - Simply use the change of a performance metric.
 - Adding or deleting a feature cannot further improves some prediction accuracy.
 - Reach the empty or full feature set.

Three Typical Methods

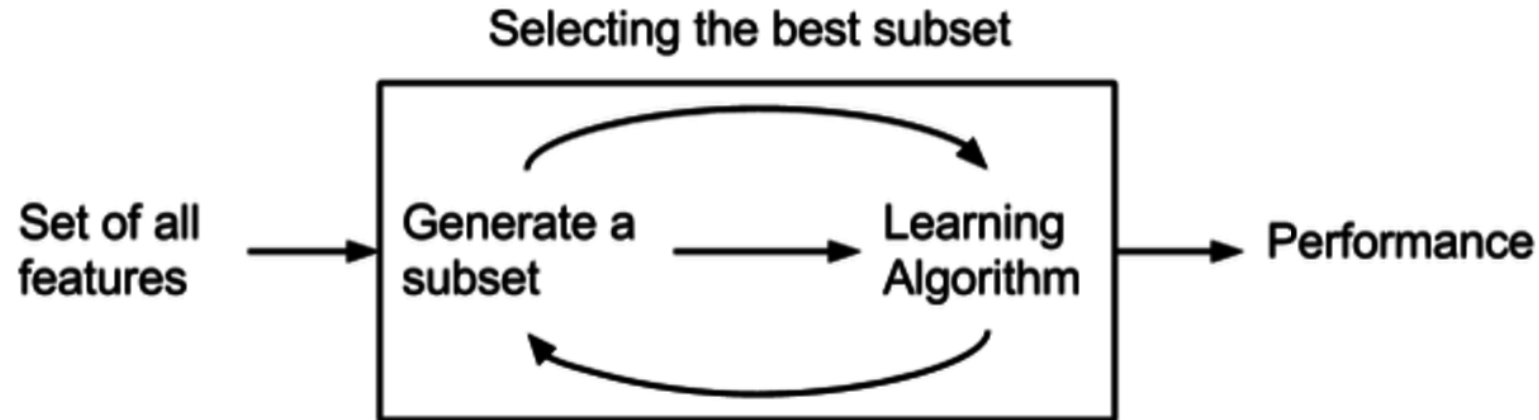
1. Wrapper methods
2. Filter methods
3. Embedded methods

1. Wrapper Methods




- **Basic idea: model dependent**
 - Navigate feature subsets by adding/removing features.
 - Evaluate the performance of the chosen model on the validation set.
 - Repeat until no improvement to the validation set accuracy.
- **Assessment: use cross validation**

1. Wrapper Methods



- **Advantage:** highly accurate
- **Disadvantage:** computationally expensive, risk of overfitting
- **Examples:** recursive feature elimination, sequential feature selection, genetic algorithms

2. Filter Methods



- **Basic idea:** independent of learning models
 - Rank features on some *heuristic* score based on their relevance to the AI task
 - Choose a subset based on the sorted scores
- **Heuristic score:** many popular scores [4]
 - Does the individual feature seem helpful in prediction?
 - Classification with categorical features: χ^2 , information gain, document frequency
 - Regression: correlation, mutual information
- **Assessment:** use statistical tests

2. Filter Methods

- Advantages
 - very fast & simple to apply
 - usually better generalization
- Disadvantage
 - not take into account interactions between features
 - not as accurate as Wrappers
- Suggestions:
 - use it as a pre-processing for further Wrapper feature selection
- Examples: Belief, correlation-based filters, fast correlated-based filters

Example: Correlation-based Filters

- **Hypothesis:** A good feature should be highly correlated to the output but not very correlated with each other.
- **Technical questions:** for a classification problem
 1. Whether a feature is **relevant** to the class?
 2. Whether a relevant feature is **redundant** with other relevant features?

Example: Correlation Scores

Two groups of correlation metrics between random variables X and Y :

1) Classical linear correlation: e.g. Pearson correlation

$$\rho(X, Y) = \frac{\sum_i (x_i - \bar{x}) (y_i - \bar{y})}{\left[\sum_i (x_i - \bar{x})^2 \cdot \sum_i (y_i - \bar{y})^2 \right]^{\frac{1}{2}}} \in [-1, +1]$$

- **Advantage:** easy and fast to compute
- **Disadvantage:**
 - cannot capture nonlinear correlation
 - calculation requires all features contain numerical values

Example: Information Gain

Two groups of correlation metrics between random variables X and Y :

2) Information theory: e.g. information gain [Quinlan, 1993]

$$IG(X; Y) = \mathcal{H}(X) - \mathcal{H}(X|Y)$$

- $\mathcal{H}(X) \triangleq -\sum_k p(x_k) \log_2 p(x_k)$ is entropy of X
- $\mathcal{H}(X|Y) \triangleq \sum_j p(Y = y_j) \cdot \mathcal{H}(X|Y = y_j)$ is conditional entropy
- Advantage: capture nonlinear correlation
- Disadvantage:
 - higher computational cost
 - IG is biased in favor of features with more values

Example: Symmetric Uncertainty

Two groups of correlation metrics between random variables X and Y :

2) Information theory: e.g. symmetric uncertainty [Press et al., 1988]

$$IG(X; Y) = \mathcal{H}(X) - \mathcal{H}(X|Y) = - \sum_j \sum_k p(x_k, y_j) \log_2 \frac{p(x_k, y_j)}{p(x_k)p(y_j)}$$

$$IG(Y; X) = \mathcal{H}(Y) - \mathcal{H}(Y|X) = - \sum_k \sum_j p(y_j, x_k) \log_2 \frac{p(y_j, x_k)}{p(y_j)p(x_k)}$$

$$SU(X; Y) = 2 \left[\frac{IG(X; Y)}{\mathcal{H}(X) + \mathcal{H}(Y)} \right] \in [0, 1]$$

- Advantage:
 - compensate for IG's bias towards features with more values
 - normalize its values to $[0, 1]$


Example: Correlation-based Filters

- Main Procedure

- 1) *C-correlation*: Use some *correlation score* to rank features according to their correlation to the class.
- 2) Ranking cut-off is determined by the user to form the *relevant feature set*.
- 3) *F-correlation*: Some relevant features are removed by redundancy detection based on the same *correlation measure*.

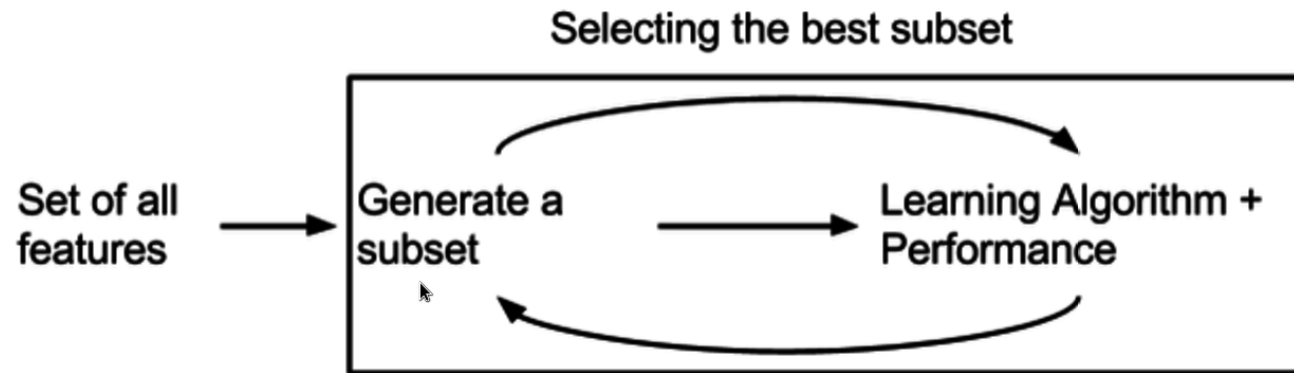
- Read paper [6] for details.

3. Embedded Methods



- **Basic idea:** Feature selection is part of model construction, and feature search is guided by the learning process.
- **Assessment:** use cross validation
- They use the specific structure of the model returned by the algorithm to get the set of 'relevant' features.

3. Embedded Methods



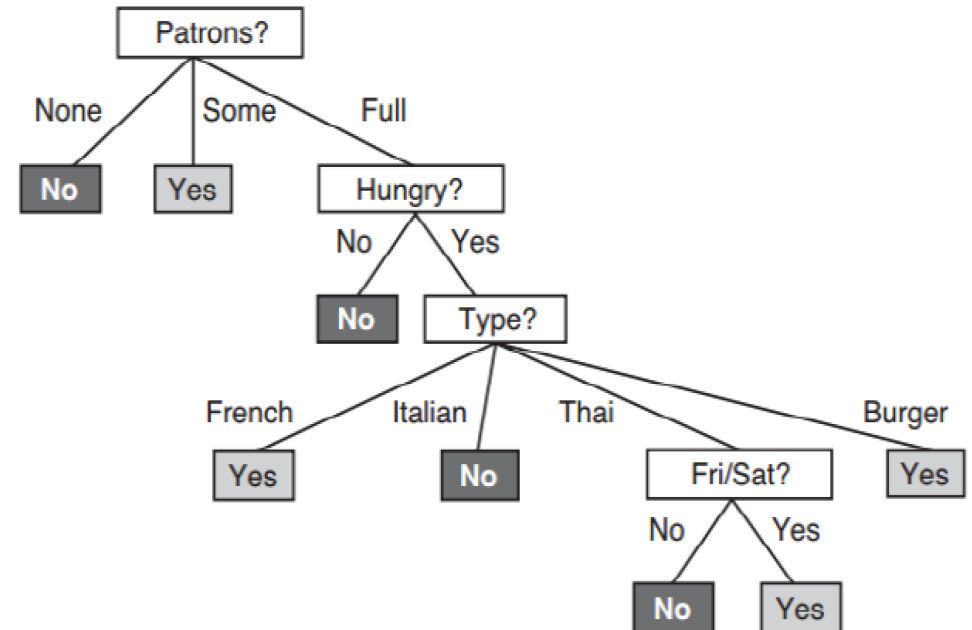
- **Advantages:**
 - similar to Wrappers, but
 - less computationally expensive & less prone to overfitting
- **Examples:** classification and regression trees, C4.5, random forest

3. Embedded Methods

- They are not too far away from wrapper techniques.
- They are a good inspiration to design new feature selection techniques for your own algorithms.
 - Find a function of features that represents your prior knowledge about what a good model is.

Example: Decision Tree

- **Review:** construct decision tree
 - Start from an empty tree.
 - Split the next best feature based on *information gain*.
 - Repeat.
- Tree construction is the process of feature selection.
- Not all features are used in the constructed tree.



Four features out of total 10 are used in constructing the decision tree.

Summary: Three Typical Methods

- Wrapper methods: model specific
- Filter methods: independent of model
- Embedded methods: feature selection is embedded in model learning

Tackling Feature Explosion

- ❖ Introduction to Feature Explosion
- ❖ Feature Selection
- ❖ Regularization (more detailed)

Regularization

- Basic idea:

- The more features matter in the model, the bigger complexity.
- Regularization = introducing **penalty** for complexity → reduce features

- Interpretation:

- It bias the model toward lower complexity (fewer features).
- Application of Occam's razor: the model should be simple (fewer coefficients).
- Bayesian viewpoint: regularization = imposing prior knowledge that the world is simple on the learning model.

Regularization Formulation

- Find $f \in \mathcal{F}$ minimizing

$$\frac{1}{n} \sum_{i=1}^n \mathcal{L}_{tr}(y_i, f(x_i)) + \lambda \cdot \Omega(f)$$

- \mathcal{F} : a class of candidate functions
- $\Omega(f)$: the complexity of a model f
- $\lambda > 0$: a regularization parameter
- **Question:** How do we pick parameter λ ?
- **Answer:** Cross validation.

Examples of Regularization Methods

- Ridge regression [Hoerl and Kennard 1970]
- Lasso regression [Tibshirani 1996]
- Smoothing splines [Wahba 1990]
- Support vector machines [Vapnik 1998]
- Regularized neural networks
- etc.

Review: Multivariate Linear Regression

- **Given:** data $\mathbf{X} \in \mathbb{R}^{N \times D}$, and output $\mathbf{y} \in \mathbb{R}^{N \times 1}$.
 - N : #samples, D : #features
- **Aim:** find $\boldsymbol{\theta} \in \mathbb{R}^{D \times 1}$ to minimize $\frac{1}{2} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2$.
- **Solution:** $\boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

Feature Selection in MLR Model

- In a MLR model, each θ_i corresponds to one feature.
- Feature selection can be treated as the penalty on θ :

$$\Omega(f) := ||\theta||_p$$


- $\theta_i = 0$: remove the i^{th} feature from the model
- ℓ_p norm of θ

Penalty ℓ_p

- **Euclidean** $p = 2$, $||\boldsymbol{\theta}||_2 = \sqrt{(\theta_1^2) + \dots + \theta_D^2}$
- ℓ_2 can be viewed as a Gaussian prior on model parameters

- **Manhattan** $p = 1$, $||\boldsymbol{\theta}||_1 = |\theta_1| + \dots + |\theta_D|$
 - ℓ_1 can be viewed as a Laplace prior on model parameters

- Generally $0 < p < \infty$, $||\boldsymbol{\theta}||_p = \sqrt[p]{|\theta_1|^p + \dots + |\theta_D|^p}$



Ridge Regression

- Ridge regression model: $\min_{\boldsymbol{\theta}} \frac{1}{2} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_2^2$
- Solution: $\boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$
- Lead to a solution with many small $\boldsymbol{\theta}$.
 - ℓ_2 does not strongly zero parameters (remove features), but still limits model complexity and get fewer features.
 - It also solves the problem that $\mathbf{X}^T \mathbf{X}$ is not invertible

Lasso Regression

- Lasso regression model: $\min_{\boldsymbol{\theta}} \frac{1}{2} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_1$
- Solution: no analytical solution
 - Need sub-gradient of ℓ_1 norm
- Lead to a sparse solution, i.e., $\boldsymbol{\theta}$ has many zero elements.
 - Remove many features and preferable for high-dimensional problems

Regression with Penalty $\ell_{1/2}$

- **Penalty $\ell_{1/2}$ model:** $\min_{\boldsymbol{\theta}} \frac{1}{2} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_{\frac{1}{2}}$
- **Solution:** non-convex and thus hard to optimize
 - Initialize with ℓ_1 penalty solution
 - Further perform gradient steps
 - Not optimal but give sparser solutions than ℓ_1 .
- Lead to an even sparser solution, and often better performance.

Remarks on ℓ_1

Two types of ℓ_1 penalty used in regression:


- Lasso for **sparsity**

$$\hat{\boldsymbol{\theta}} = \arg \min_{\boldsymbol{\theta}} \frac{1}{2} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_1$$

- ℓ_1 loss for **robustness**

$$\hat{\boldsymbol{\theta}} = \arg \min_{\boldsymbol{\theta}} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_1 + \lambda \|\boldsymbol{\theta}\|_p$$

Remarks on ℓ_1 Continue



ℓ_1 penalizes more than ℓ_2 when θ is small \rightarrow use for sparsity

ℓ_1 penalizes less than ℓ_2 when θ is big \rightarrow use for robustness

Summary

- Feature engineering is often crucial to get good results.
- Manual feature learning requires knowledge for the task.
- Automated feature learning is much more preferred.
- Strategies for tackling feature explosion:
 - Feature selection is a *heuristic* search problem.
 - Use regularization on all possible features to prevent overfitting.

Reading Materials For This Lecture

- [1] Online course: <http://clopin.net.com/isabelle/Projects/ETH/>
- [2] A. Zheng. and A. Casari. 2017. *Mastering Feature Engineering for Machine Learning Models*. Chapter 3.
- [3] Curse of dimensionality: https://en.wikipedia.org/wiki/Curse_of_dimensionality
- [4] Y. Yang and J. O. Pedersen. 1997. *A Comparative Study on Feature Selection in Text Categorization*. ICML. pp:412-420.
- [5] Blog: <https://machinelearningmastery.com/an-introduction-to-feature-selection/>
- [6] L. Yu and H. Liu. 2003. Feature Selection for High-dimensional Data: A Fast Correlation-based Filter Solution. ICML. pp: 856-863
- [7] An Introduction to Feature Extraction. Isabelle Guyon and André Elisseeff.
- [8] <https://developers.google.com/machine-learning/data-prep>